

Calculation of cache misses

Variant 0:

n: The number of columns and number of rows in the matrix. (n x n matrix)

ecl: Elements per cache line

entry: An element of the matrix

cl: A single cache line

$$\text{ecl} = \text{sizeof}(\text{entry}) / \text{sizeof}(\text{cl}) = 32 / (64 * 8) = 16$$

$$\text{Misses} = (3/\text{ecl}) * n^2 = (3/16) * n^2$$

$$\text{Accesses} = 3 * n^2$$

$$\text{Misses in \%} = \text{Misses} / \text{Accesses} = ((3/\text{ecl}) * n^2) / (3 * n^2) = 1/\text{ecl} \\ = 1/16 = 6.6\%$$

Variant 1:

Under the Assumption that a complete Row of the Matrix cannot be stored in Cache. Otherwise on the next outer Loop Iteration the Cache would be able to return the first Column.

$$\text{ecl} = \text{sizeof}(\text{element}) / \text{sizeof}(\text{CacheLine}) = 32 / (64 * 8) = 16$$

$$\text{Misses} = 3 * n^2$$

$$\text{Accesses} = 3 * n^2$$

$$\text{Misses in \%} = \text{Misses} / \text{Accesses} = (3 * n^2) / (3 * n^2) = 100\%$$

Findings:

When we go through the matrix in row major order we get a missrate of 3.3% on write operations and 0.7% on read operations. If we iterate with column major through the matrixes we notice a cache miss rate of 99.9% on writes but only 14% in reads in the level 1 cache on cachegrind. Cachegrind provides us with a simulated environment and perf reports the actual hardware performance. Therefore it could result in the difference to cachegrind. Here we found a cache miss rate of 0.231% for row major and 1.909% for column major. We assume there is a special prediction in the caches and therefore we get a lot less misses as we expected.